

How GenAI Systems Actually Work

The diagrams behind the timeline — now organized by phase. Every concept that's easier to see than to read, from how a model is born to how you evaluate, ground, bound, and ship one, drawn the way a senior would sketch it on a whiteboard to make it click.

COMPANION TO → "From LLM Integration to Production Agentic Systems" · 19 diagrams across 5 phase bands · Color is consistent: blue = model · teal = retrieval · amber = data/eval · rose = action/safety · violet = orchestration · slate = infrastructure

THE DIAGRAMS – ORGANIZED BY PHASE

01 How a model is born: pretraining → fine-tuning → RLHF

02 From text to tokens to numbers

03 What an LLM does in one forward pass

04 Inside the transformer: attention

05 How the next token is chosen

06 The context window & why memory is an illusion

07 Anatomy of a prompt + chain-of-thought

08 The evaluation system (the interview-winner)

09 LLM-as-judge & its biases

10 The RAG pipeline, end to end

11 Embeddings & vector space

12 Chunking strategies compared

13 Hybrid search & re-ranking

14 RAG eval triad: which stage failed?

15 Query transformation & HyDE

16 The agent loop

17 Multi-agent orchestration

18 Prompt injection & guardrails

19 Production architecture + eval-gated CI/CD

Closing reference → the full capability stack, mapped to phases


How the model itself works

Before any application, you must be able to explain the machine. These six cover what an LLM *is* mechanically — how it's trained, how it reads text, how one forward pass produces a token, how attention works, how sampling picks, and the finite context window everything must fit inside. This is the layer interviewers probe to see if you understand the engine or just call the API.

DIAGRAM 01 · FOUNDATIONS

How a model is born: pretraining → fine-tuning → RLHF

Before you can reason about an LLM's behavior, you need to know where it came from. A model passes through three stages, each shaping it differently. Understanding this explains *why* models behave as they do — why they're fluent but not automatically helpful, why they have a knowledge cutoff, and why "the model was trained on X" is a different claim from "the model can retrieve X." This is also the track this curriculum deliberately does *not* send you down — you build systems *around* trained models — but you must be able to discuss it.



TRAINING (ONCE, OFFLINE) VS INFERENCE (EVERY CALL)


The crucial split for your role: all three stages happen *before* you ever touch the model. When you "call the LLM" you're doing *inference* on frozen weights — you change its behavior through prompts and retrieved context, not by retraining. That's why RAG exists: it's how you give a frozen model fresh, private knowledge without going back to Stage 1.

Interview translation: "fine-tuning" and "RAG" solve different problems and interviewers test if you conflate them. Fine-tuning (Stage 2) changes *behavior/style* and needs retraining; RAG injects *knowledge* at query time with no retraining. "Should we fine-tune or use RAG?" → usually RAG for knowledge, fine-tuning for form. Knowing the difference cleanly is a senior signal.

DIAGRAM 02 • FOUNDATIONS


From text to tokens to numbers

An LLM never sees letters or words. The very first thing that happens to your prompt is **tokenization** — text is chopped into sub-word fragments, and each fragment is swapped for an integer ID. Those IDs are then looked up in an **embedding table** to become vectors. Everything downstream is math on numbers; the language is gone almost immediately.



What an LLM does in one forward pass


Strip away the mystique and an LLM is a **next-token predictor**. It takes the whole sequence of token-vectors, runs them through a deep stack of transformer layers, and emits a probability distribution over the entire vocabulary for *what comes next*. It picks one token, appends it, and runs the whole thing again. Generation is this loop, one token at a time — which is exactly why latency scales with output length.



Interview translation: when you say "the model hallucinated," you're describing this picture — it sampled a high-probability-but-wrong token. There's no truth-check anywhere in the pass. Constraining that is the entire job of RAG, citations, and validation.

Inside the transformer: the attention mechanism

The 2017 idea that made all of this work. When the model processes one token, **attention** lets it look at every other token in the context and decide how much each one matters *right now*. That's how it resolves "it" back to the noun it refers to, or connects a question to its answer across a paragraph. Each token builds a new, context-aware representation of itself by pulling in a weighted blend of all the others.



How the next token is actually chosen

The model hands you a probability distribution; **sampling parameters** decide how you pick from it. This is the knob that controls whether your system is repeatable (essential for evals and tool-calling) or varied (good for drafting). Understanding it answers the classic "why is my output non-deterministic?"

temperature = 0

always take the top token · deterministic

mat 0.61 ✓ PICKED

floor 0.18

rug 0.09

sofa 0.05

Same prompt → same answer,
every single time.

Use for: evals, tool-calling,
anything needing reproducibility.

temperature = 0.9

flatten the odds · sample · varied

mat 0.40

floor 0.27 ✓

rug 0.20

sofa 0.13

Probabilities are softened, so
lower-ranked tokens can win.

Use for: brainstorming,
creative drafting, variety.

top-p = 0.9 (nucleus)

keep smallest set summing to 0.9

mat 0.61 ← in

floor 0.18 ← in

rug 0.09 ← in

sofa 0.05 ✗ cut

Chop off the unlikely tail first,
then sample from what remains.


A cleaner randomness knob —
often tuned alongside temperature.

THREE WAYS TO PICK

Greedy (temp 0) is a special case: "always pick the most likely." Temperature reshapes the whole distribution; top-p truncates its tail. Knowing the difference is a small but reliable senior tell.

The context window — why "memory" is an illusion

An LLM has **no memory** between calls. Its only working space is the **context window**: a fixed token budget holding everything it can currently see. Every "memory" feature you'll build — chat history, RAG, agent scratchpads — is really just a strategy for deciding *what to load into this finite box* before each call. Internalize this and most "naive implementation" bugs explain themselves.



Controlling and measuring an unreliable component

The prompt is your only control surface over a model you can't retrain — so Phase 1 treats it as engineering, not wording tricks, and wraps it in measurement. These three diagrams cover prompt anatomy, the evaluation harness that both documents call *"the deliverable that wins interviews,"* and the LLM-as-judge technique with its biases. **This is the phase most candidates skip and the one that most separates senior from junior.**

DIAGRAM 07 • PHASE 1

Anatomy of a prompt + chain-of-thought

A production prompt is a structured object with distinct parts, each doing one job. The naive approach jams everything into one blob; the senior approach separates **system instructions**, **few-shot examples**, clearly **delimited** data, and an explicit **output contract**. **Chain-of-thought** — asking the model to reason before answering — is a separate lever that trades tokens for accuracy on hard problems.

A WELL-STRUCTURED PROMPT (each block = one job)

① System prompt

role, rules, persona — highest-priority, sets behavior for the whole call

② Few-shot examples

2–5 (input → ideal output) pairs that demonstrate the pattern.

Examples teach faster than instructions. (zero-shot = none; few-shot = a few)

③ Delimited context / data

<context> ...retrieved chunks or user data... </context>

Clear fences stop the model confusing data for instructions → injection defense

④ Output contract

"Return JSON: {category, confidence}." Exact format, ideally schema-enforced.

Vague asks → varied answers. The contract is how you get parseable output.

CHAIN-OF-THOUGHT: A SEPARATE LEVER

Without CoT

Q: $17 \times 24 + 9 = ?$

A: 417 × (guessed in one step)

Model jumps straight to an answer;
on multi-step problems it often slips.

With CoT — "think step by step"

$17 \times 24 = 408$

$408 + 9 = 417$

A: 417 ✓ (reasoned, reliable)


The model uses its own output as a
scratchpad. Costs more tokens —
senior judgment = when it's worth it.

STRUCTURE IS CONTROL

Each block is independently tunable and testable — which is exactly what makes a prompt versionable in your Phase 1 registry. CoT is drawn separately on purpose: it's not part of every prompt, it's a deliberate accuracy-for-tokens trade you reach for on reasoning-heavy tasks and measure before keeping.


The evaluation system — the interview-winner

This is the diagram to internalize above all others. An **eval harness** is a systematic, repeatable measurement of output quality. Without it you change a prompt, it "feels" better, and you've silently broken three other cases. With it, every change is scored against a **golden set**, and a regression *fails the build* before it ships. Both documents call this the single strongest senior signal — here's why, drawn out.



LLM-as-judge & its biases

Exact-match scoring works for classification, but how do you score "is this answer *helpful*?" You use a stronger LLM as a **judge**, scoring outputs against a rubric. It's powerful and scalable — but judges have known **biases**, so you **calibrate** against human labels before trusting one. Being able to name the biases is the senior tell; anyone can say "use an LLM to grade it."




Grounding the model in your own data

The most employable skill set in the curriculum. Phase 1 gave you the model and the means to measure it; Phase 2 makes it answer from *your* corpus with citations instead of hallucinating from memory. These six diagrams go deep: the full pipeline, the vector space it runs on, the chunking decisions that quietly cap quality, hybrid retrieval, the eval triad that pinpoints failures, and query transformation. **The throughline: retrieval quality is everything, and you prove it with numbers.**

DIAGRAM 10 • PHASE 2

The RAG pipeline, end to end

RAG has two halves people often blur together. **Indexing** happens *ahead of time* (offline): documents are chunked, embedded, and stored. **Retrieval + generation** happens *per query* (online): the question fetches relevant chunks, which are injected into the prompt so the model answers from real sources instead of memory. Separating these two timelines is the first thing a senior clarifies.




TWO TIMELINES, ONE STORE

The vector store is the seam where the offline and online lanes meet. The senior principle to memorize: *retrieval caps generation*. If the right chunk never enters the prompt, no amount of prompt-tuning produces a correct answer — it can only hallucinate or refuse.


Embeddings & vector space

An **embedding** turns text into a point in high-dimensional space where *distance means dissimilarity of meaning*. Texts that mean similar things land near each other even with zero shared words. Retrieval is then just geometry: embed the query, find its nearest neighbors. Here it's flattened to 2D for the eye — real embeddings have hundreds or thousands of dimensions.




Chunking strategies compared

Before text can be embedded it must be split into **chunks** — and *how* you split silently determines answer quality more than almost any other choice. Too small and a chunk loses the context that makes it meaningful; too large and retrieval becomes imprecise. This is one of the design decisions you'll A/B test and *defend with a number* in Phase 2, so you need to see the options side by side.



Hybrid search & re-ranking

Production retrieval rarely trusts one method. **Dense** (embeddings) is great at meaning; **sparse** (keyword/BM25) is great at exact terms. You run both, **fuse** their rankings (RRF), then send a small candidate set to a slow-but-accurate **cross-encoder** that re-scores query-and-chunk together. Two cheap-then-expensive stages — the pattern that lifts real RAG above tutorial RAG.




CHEAP-AND-BROAD, THEN EXPENSIVE-AND-PRECISE

The two-stage shape is deliberate: you can't run the accurate cross-encoder over a million chunks (too slow), so dense+sparse cheaply narrow to ~20, then the cross-encoder precisely orders those. This is the exact result you'll quantify in Phase 2: "hybrid + re-rank lifted recall@5 from X to Y."


The RAG eval triad — which stage failed?

When a RAG answer is bad, the killer question is *which stage broke* — retrieval or generation? The **eval triad** decomposes failure so you fix the right thing instead of blindly tuning the answer prompt. **Context precision** grades retrieval; **faithfulness** grades whether the answer stuck to the context; **answer relevance** grades whether it addressed the question. Three metrics, three different culprits.



Query transformation & HyDE

The user's raw question is often a *poor search query* — short, vague, phrased unlike the documents that answer it. **Query transformation** fixes this before retrieval. The cleverest variant is **HyDE** (Hypothetical Document Embeddings): instead of searching with the question, have the LLM *imagine an answer* and search with *that* — because a hypothetical answer sits closer in vector space to the real answer than the question does.




Making the system act — reliably

Now the system takes actions, not just answers. The senior inversion: the skill is *reliability and control*, not autonomy. These two diagrams cover the bounded single-agent loop and multi-agent orchestration — including the compounding-error math that is the strongest argument *against* reaching for multiple agents. The RAG service from Phase 2 becomes a *tool* the agent calls.

DIAGRAM 16 • PHASE 3

The agent loop

An **agent** is an LLM placed inside a loop where it can take actions. The cycle is **reason** → **act** → **observe**, repeating until the task is done or a limit trips. Demystified: it's a **while** loop around tool-calling. The senior emphasis isn't autonomy — it's the *guardrails* drawn in rose: budgets, validation, and a human gate before anything irreversible.




REASON • ACT • OBSERVE, BOUNDED

The lock icons mark irreversible actions that must pass a human gate — the agent may *propose* freely but only *commit* through approval. "A flashy agent that fails 30% is a toy; a bounded one that fails 2% and degrades safely is a senior artifact."

Multi-agent orchestration

When a task genuinely decomposes, a **supervisor** coordinates specialist agents instead of one generalist juggling everything. Useful for truly separable subtasks — and expensive theater otherwise. The diagram shows the common supervisor/worker topology, plus the failure mode that is the strongest argument *against* reaching for it: compounding error.



SUPERVISOR / WORKER TOPOLOGY

The clearest senior signal in this whole area is being able to argue *against* your own multi-agent design. The compounding-error math is why: chaining unreliable steps multiplies the failure. Reach for multi-agent only when subtasks are genuinely separable and the quality gain beats the cost and reliability hit.


Shipping it — hardened and observable

The phase where your existing AWS and backend seniority becomes a genuine edge — most GenAI candidates can't operate what they build. These two diagrams cover the safety layer (especially the RAG-borne prompt injection that catches people off guard) and the full production architecture with the eval-gated deploy pipeline that ties the whole curriculum together.

DIAGRAM 18 • PHASE 5


Prompt injection & guardrails

Prompt injection is the defining security problem of LLM apps: malicious text that hijacks the model's instructions. The obvious version arrives in user input — but the dangerous one rides in on a *retrieved document* your RAG system trusts. Because the model can't reliably tell *instructions* from *data*, a poisoned chunk can carry commands. Knowing the RAG-borne vector specifically is a senior-level awareness.




Production architecture + eval-gated CI/CD

The capstone picture: your system running on AWS, instrumented, guard-railed, and — the senior centerpiece — deployed through a pipeline that **runs your eval suite and blocks the deploy if quality regresses**. This is where the whole curriculum converges: the eval harness from Phase 1 becomes a release gate, exactly like tests gate normal code.



The whole stack, one picture

Finally, how the layers compose. Generative AI is the foundation; the LLM is the engine; RAG grounds it in your data; an agent wraps it in a decision loop; agentic orchestration coordinates several. Each layer *consumes the one below and adds capability* — and your Phase artifacts map directly onto these layers.



How to use these 19 diagrams: don't memorize them as pictures — be able to *redraw* each one on a whiteboard while narrating it. Work through them by phase band: foundations first (can you explain the engine?), then prompting/eval, then RAG, then agents, then production. An interviewer who watches you sketch the eval loop or the RAG triage triad, explaining each arrow, learns more about your seniority in ninety seconds than a résumé conveys in a page. The diagram that does the most work for you is **08 (evaluation)** — if you can draw and defend that one cold, you've cleared the bar most candidates trip on.

VISUAL COMPANION · to "From LLM Integration to Production Agentic Systems"

Color legend — ■ model/LLM · ■ retrieval/RAG · ■ data/eval · ■ action/safety · ■ orchestration · ■ infrastructure. 19 diagrams across 5 phase bands, every one hand-built SVG — open this file in any browser, zoom freely, nothing pixelates.